Instance Segmentation with Mask R-CNN

AI Computer Vision Project, AI Innovation Square

2020.03.23 – 2020.04.10

JIHUN KIM
YUJIN NAM
YOONSUNG LEE

AI Innovation Square

1. Mask R-CNN

1-1. What's Mask R-CNN?
1-2. Problem

2. Modified Mask R-CNN

2-1. What's the difference with original Mask R-CNN?
2-2. Specified Architecture

3. Conclusion

1-1. What's Mask R-CNN?

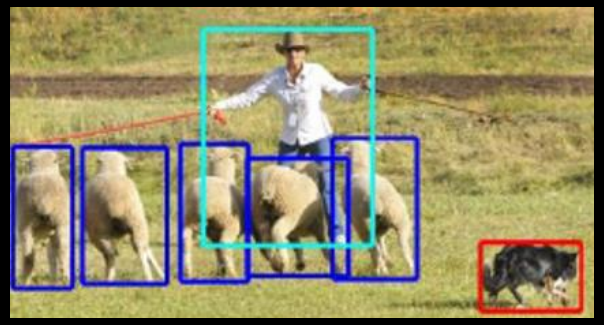

BBOX
Classification

Can Separate
Cannot Segment

Faster R-CNN

+

Segmentation
Classification

Cannot Separate
Can Segment

FCN

=

Segmentation
In Bbox
Classification

Head Architecture

Faster R-CNN + Binary Mask Prediction + FCN + RoI Align

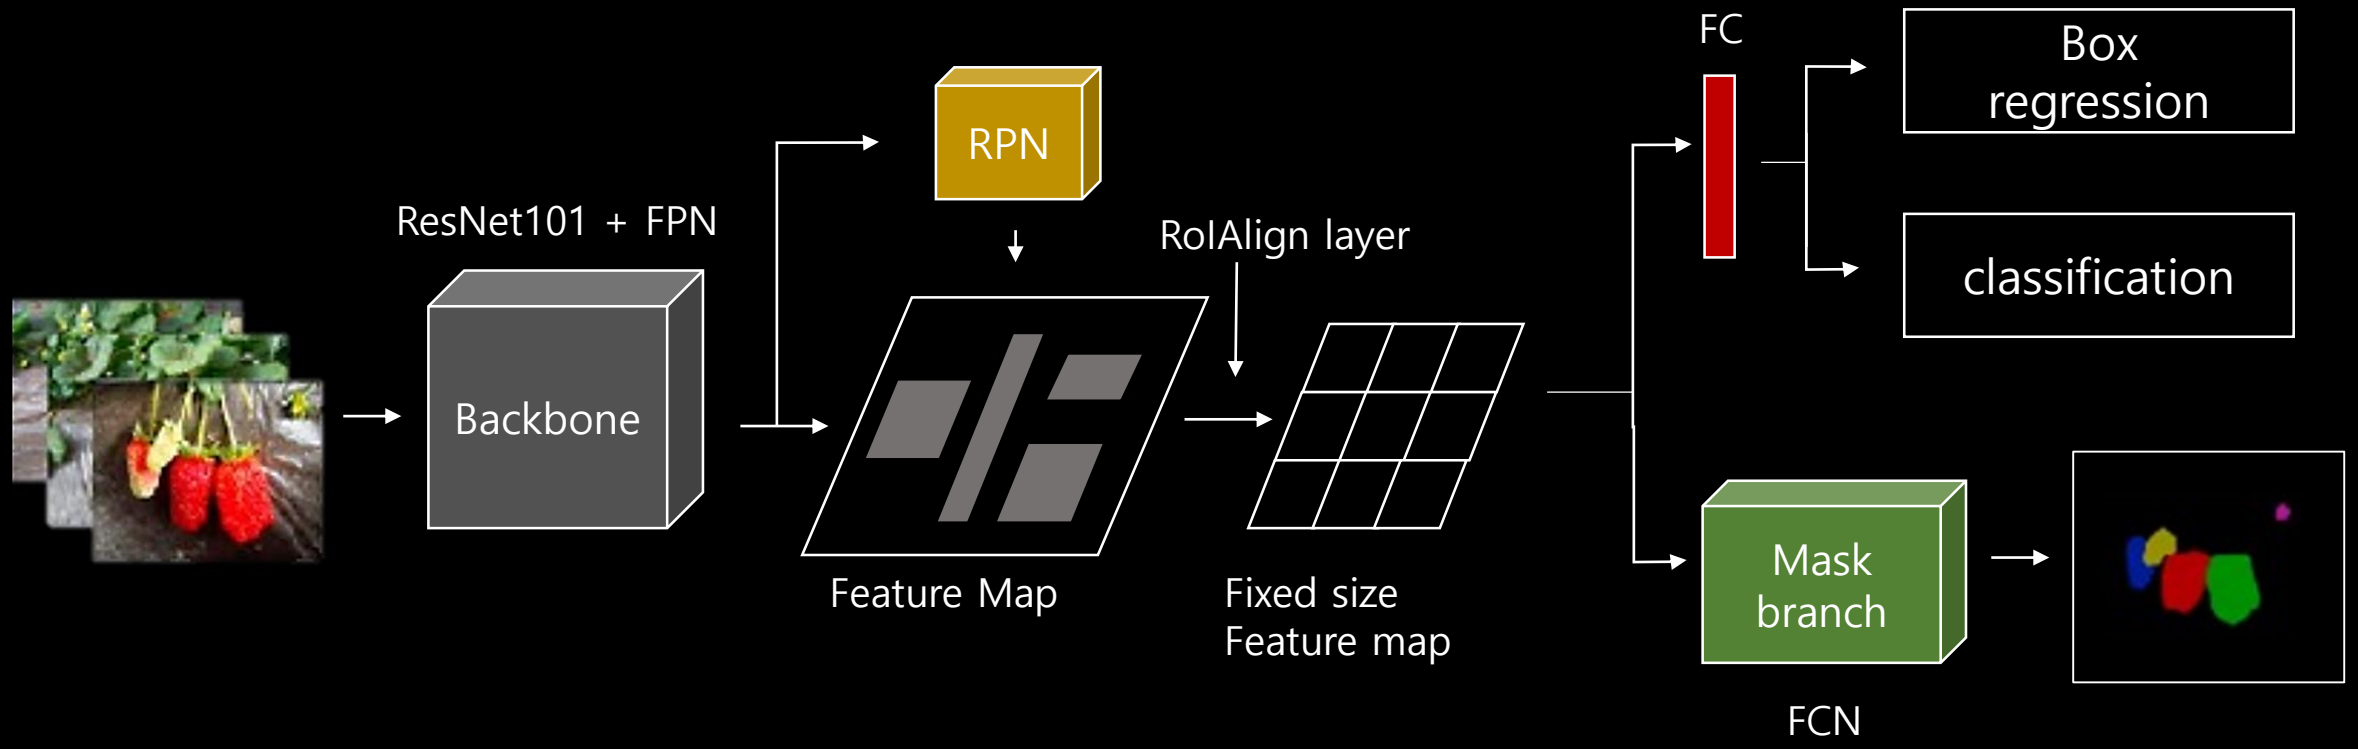

How to use Mask R-CNN?

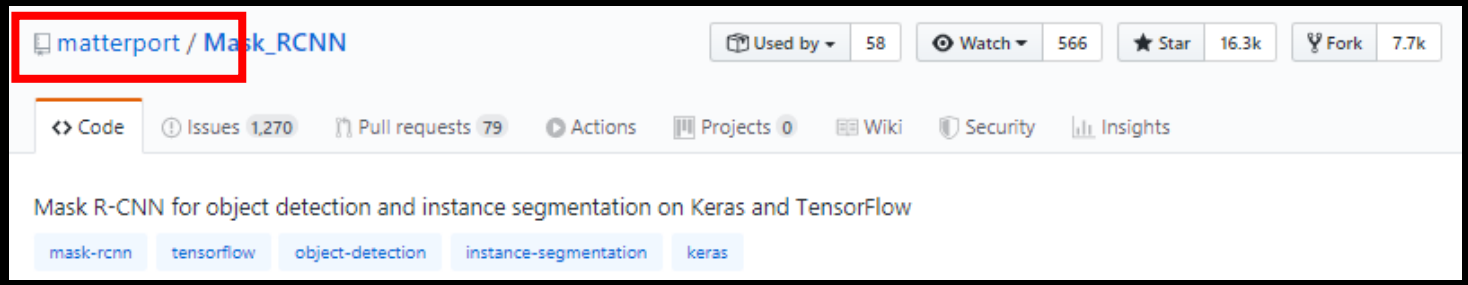

ResNet101(backbone) + FPN(Binary Masking)

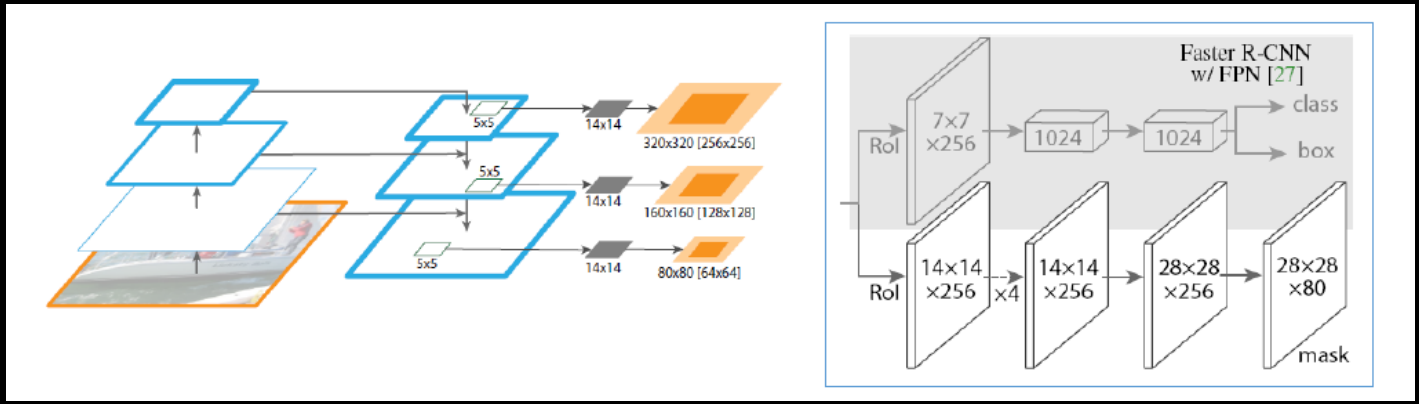

COCO Dataset

COCO
Common Objects in Context

How it works?


person 0.996
person 0.863
cell phone 0.942
cell phone 0.803
cell phone 0.990
cell phone 0.947
handbag 0.869

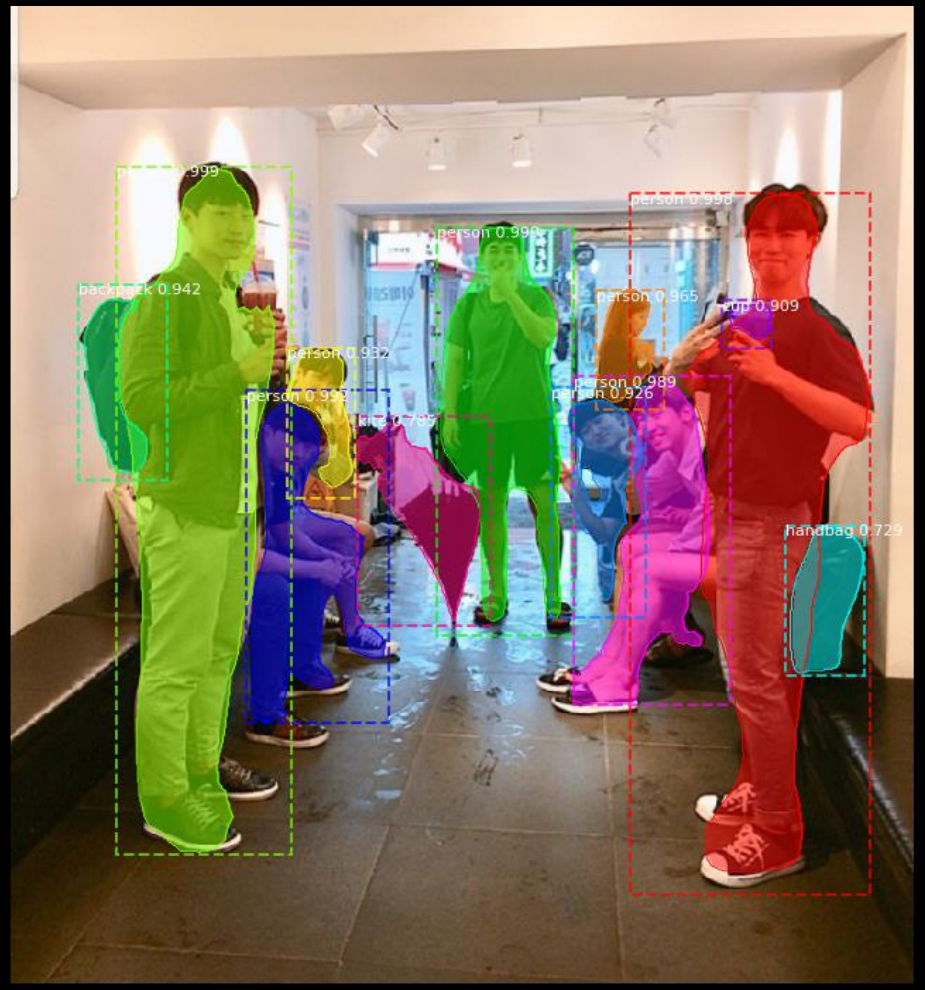
person 0.990
backpack 0.942
person 0.965
cell 0.909
person 0.989
person 0.926
handbag 0.729

How it works?

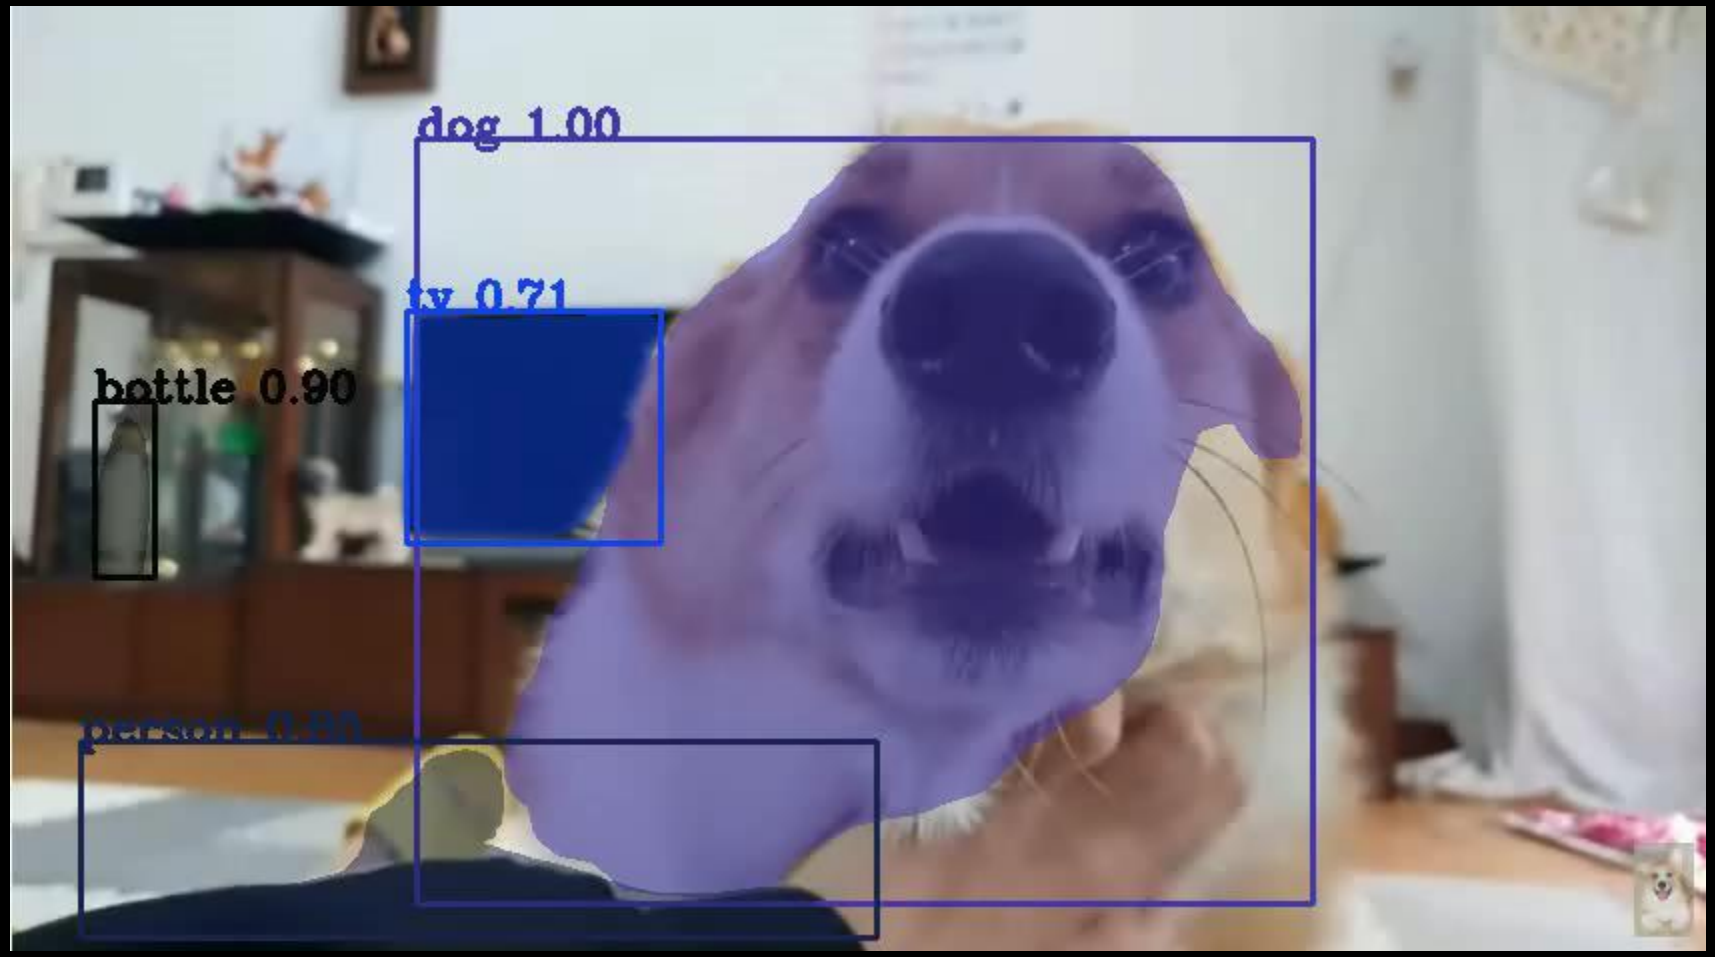

1-2. Problem


person 0.996
person 0.863
cell phone 0.990
cell phone 0.947
handbag 0.969

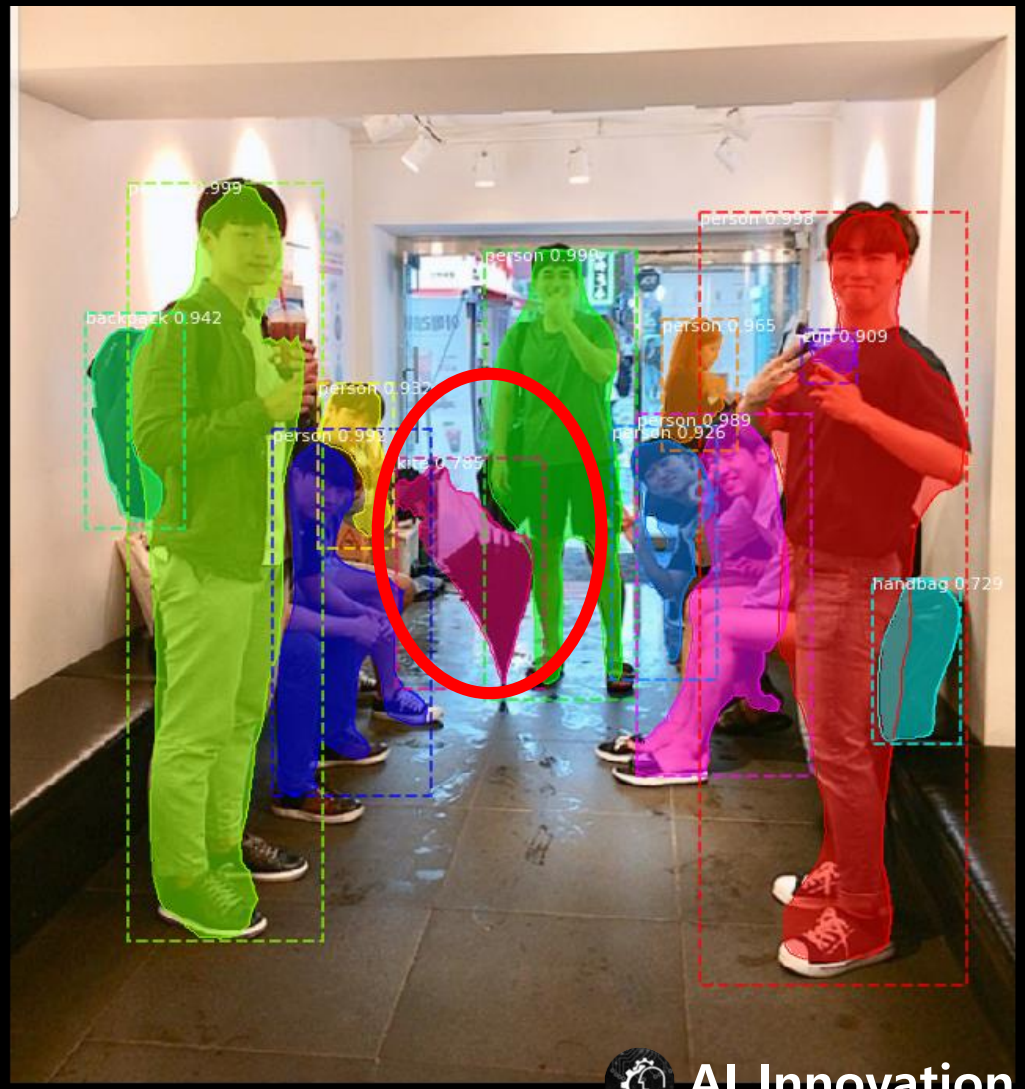
person 0.990
backpack 0.942
person 0.965
person 0.989
person 0.926
handbag 0.729

2-1. What's the difference with original Mask R-CNN?

Head Architecture

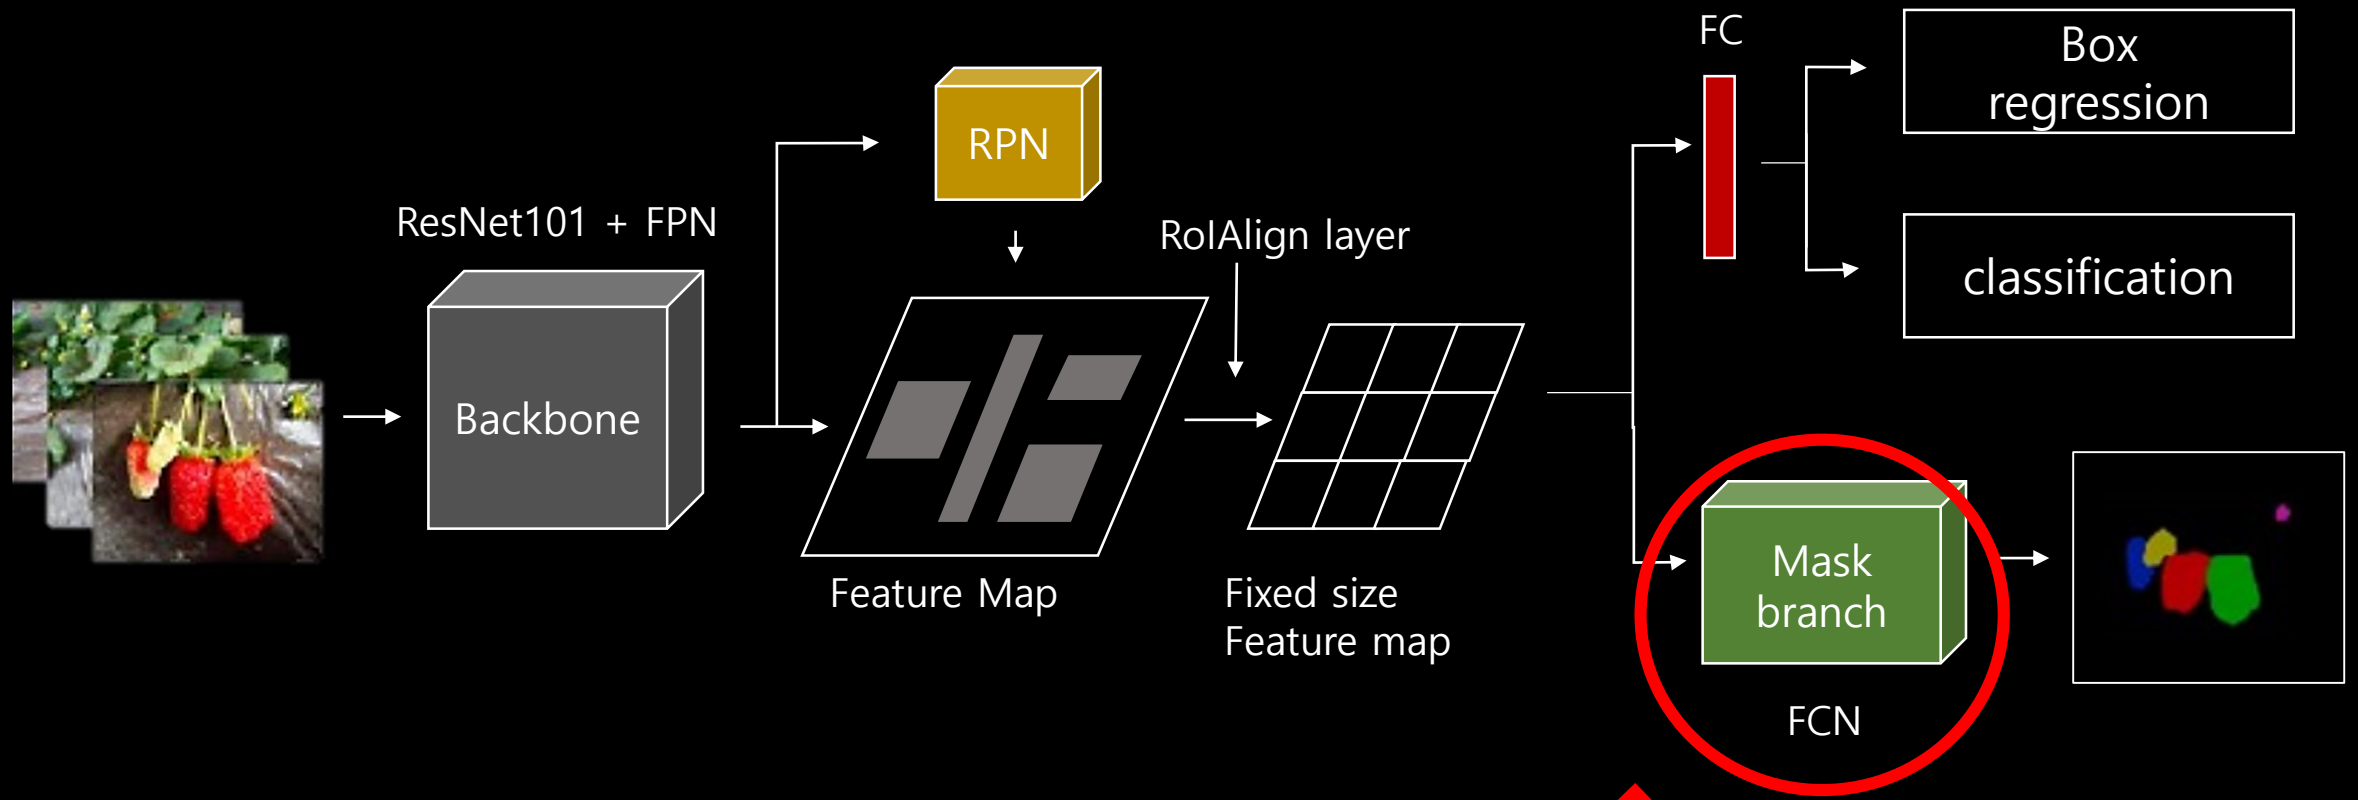

U-Net

2-1. What's the difference with original Mask R-CNN?

Modified U-Net

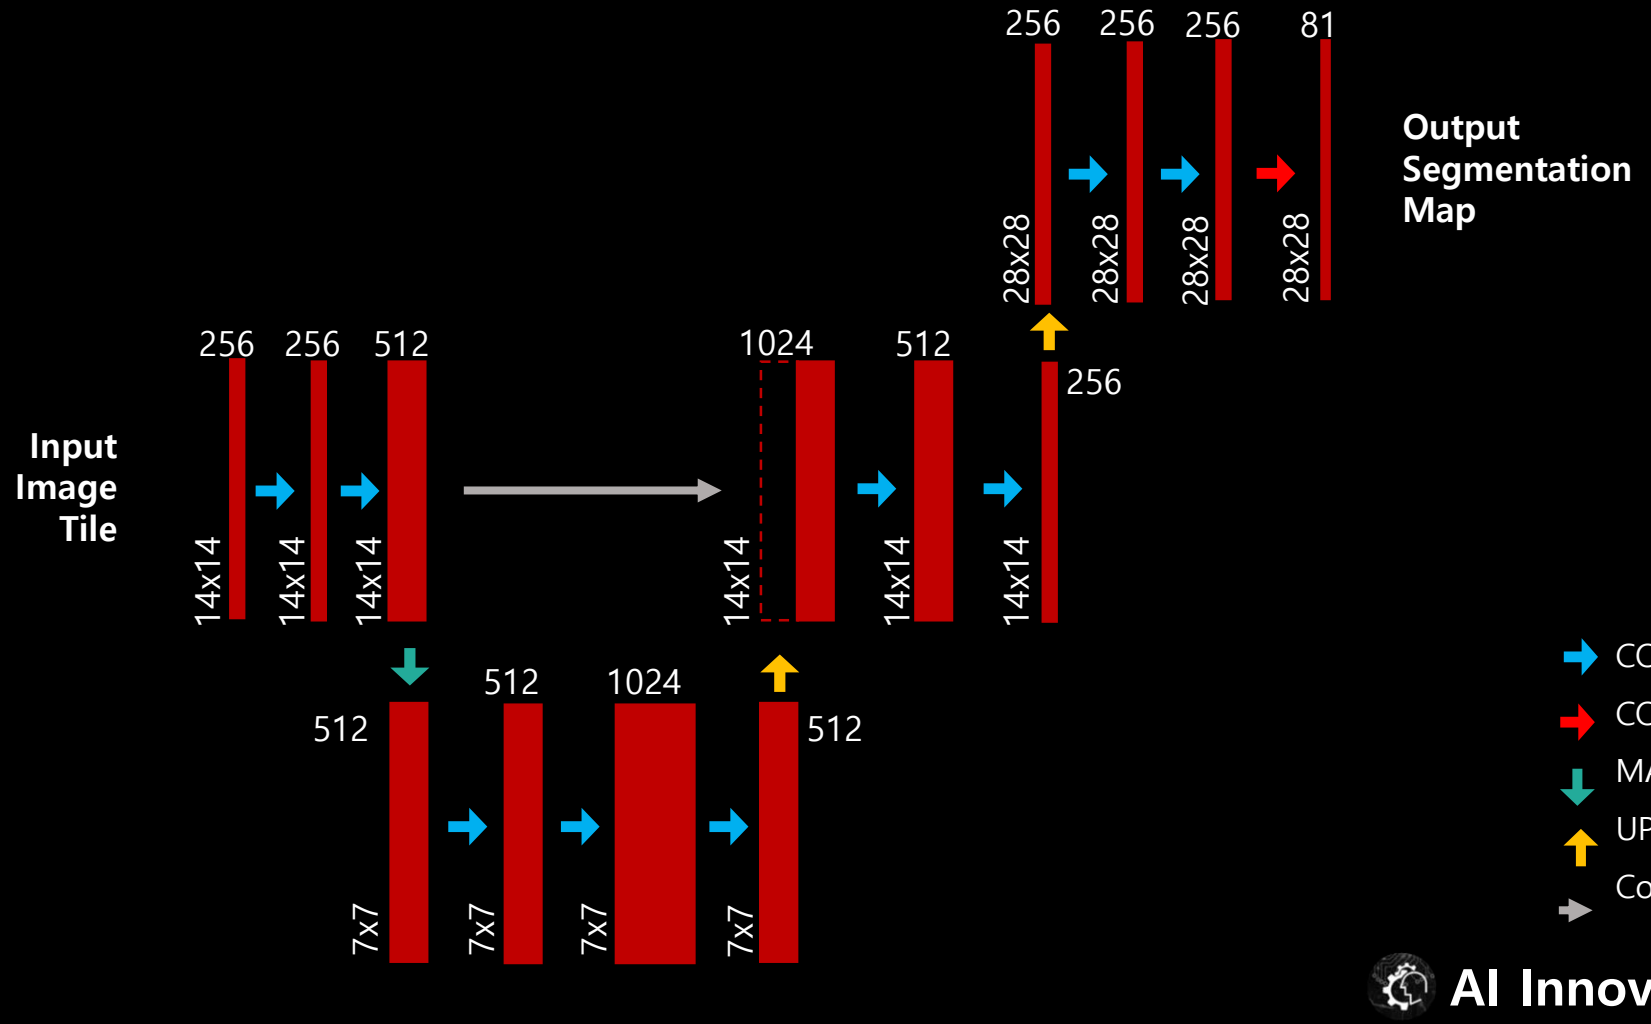

AI Innovation Square

2-1. What's the difference with original Mask R-CNN?

```
def build_unet_mask_graph(rois, feature_maps, image_meta,

                          pool_size, num_classes, train_bn=True):
    # ROI Pooling
    # Shape: [batch, num_rois, MASK_POOL_SIZE, MASK_POOL_SIZE, channels]
    # num_rois: number of regions of interest to be used
    x = PyramidROIAlign([pool_size, pool_size],
                        name="roi_align_mask")([rois, image_meta] + feature_maps)

    # Conv layers
    # (1, 100, 14, 14, 256)

    # Downsampling
    x = KL.Conv2D(256, (3, 3), padding='same', name='layer11', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNorm()(x)
    # (1, 100, 14, 14, 256)
    x = KL.Conv2D(256, (3, 3), padding='same', name='layer12', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNorm()(x)
    # (1, 100, 14, 14, 256)
    skip_connection = x  # for skip connection
    x = KL.Maxpooling2D()(x)
    # (1, 100, 7, 7, 256)


    # Bottleneck
    x = KL.Conv2D(256, (3, 3), padding='same', name='layer21', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNorm()(x)
    # (1, 100, 7, 7, 256)
    x = KL.Conv2D(512, (3, 3), padding='same', name='layer22', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNorm()(x)
    # (1, 100, 7, 7, 512)
    x = KL.Conv2D(256, (3, 3), padding='same', name='layer23', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNorm()(x)
    # (1, 100, 7, 7, 256)
```

2-1. What's the difference with original Mask R-CNN?

```
# Upsampling
x = KL.UpSampling2D()(x)
# (1, 100, 14, 14, 256)
x = KL.Concatenate(axis=-1)([x, skip_connection])
# (1, 100, 14, 14, 512)
x = KL.Conv2D(256, (3, 3), padding='same', name='layer31', activation='relu', kernel_initializer='he_normal')(x)
x = BatchNorm()(x)
# (1, 100, 14, 14, 256)
x = KL.Conv2D(128, (3, 3), padding='same', name='layer32', activation='relu', kernel_initializer='he_normal')(x)
x = BatchNorm()(x)
# (1, 100, 14, 14, 128)
x = KL.UpSampling2D()(x)
# (1, 100, 28, 28, 256)
x = KL.Conv2D(256, (3, 3), padding='same', name='layer41', activation='relu', kernel_initializer='he_normal')(x)
x = BatchNorm()(x)
# (1, 100, 28, 28, 256)
x = KL.Conv2D(256, (3, 3), padding='same', name='layer42', activation='relu', kernel_initializer='he_normal')(x)
x = BatchNorm()(x)
# (1, 100, 28, 28, 256)
x = KL.Conv2D(81, (1, 1), padding='same', name='layer43', activation='sigmoid')(x)
# (1, 100, 28, 28, 81)


return x
```

2-2. Specified Architecture

Modified Head Architecture

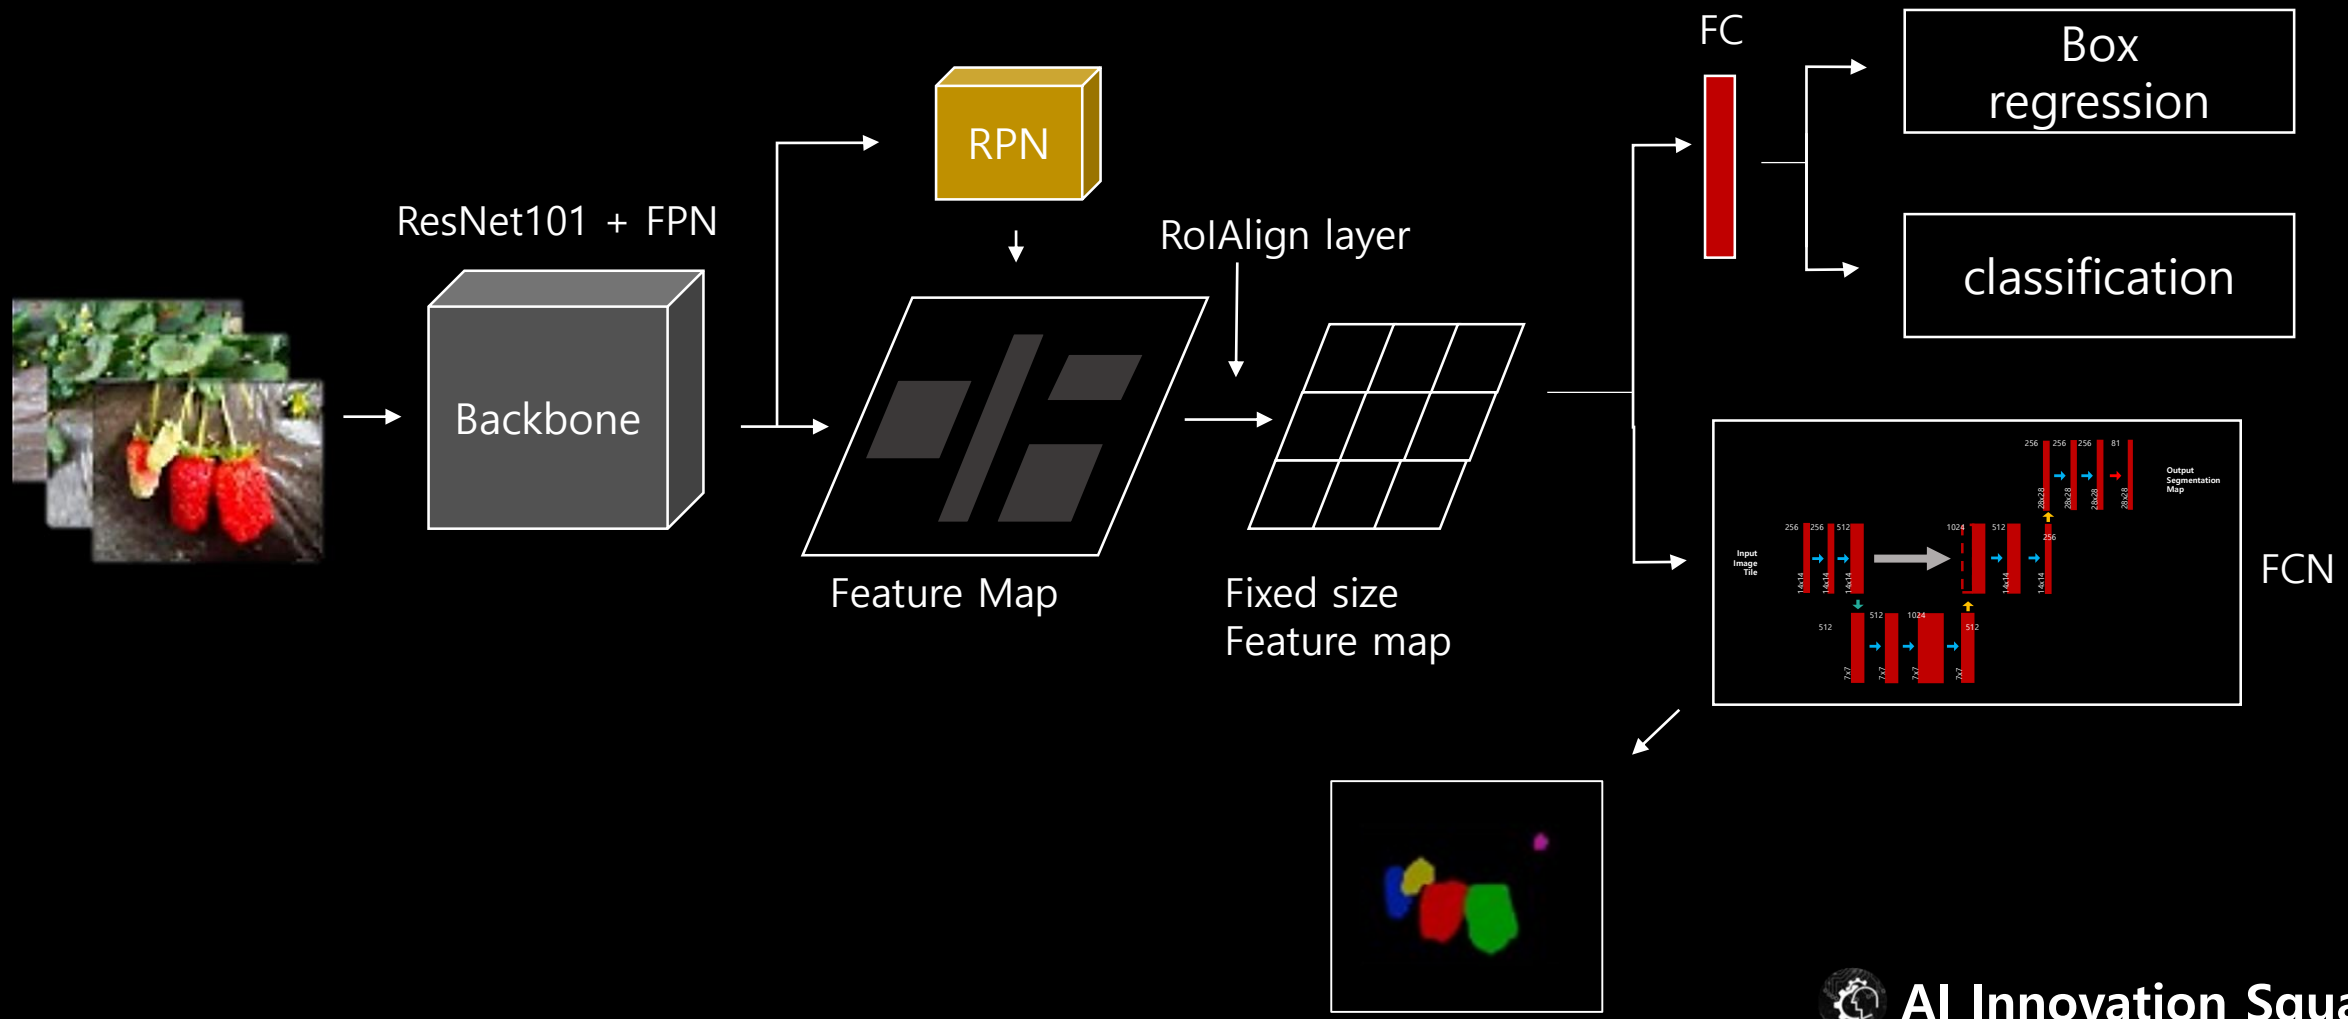

3. Conclusion

Paper Review

↓

Build an Architecture

↓

Train CoCo dataset

↓

Transfer learning
with iMaterialist dataset

Dataset : iMaterialist (Fashion) 2019 at FGVC6 (19GB)

GPU : GeForce RTX 2080 Ti
CUDA Toolkit 9.0
Anaconda (python3.7)

---

numpy
scipy
Pillow
cython
matplotlib
scikit-image
tensorflow>=1.3.0
keras>=2.0.8
opencv-python
h5py
imgaug
IPython[all]

Q & A